

Big Data and Data Mining

Supervised Learning

Flavio Bertini

flavio.bertini@unipr.it

Regression

- The regression problem asks to predict a numerical variable's value given the values of other variables
- Easiest example is with two variables x and y :
 - x is the variable in input
 - y is the variable we want to predict

x	y
1	3
2	5
3	7
4	9
5	11

Training data: x is the input data, y is the label data

Learning

Learning task: **what's the mapping from x to y ?**

Testing

Try to "learn" from the training data an hypothesis function h

What is $h(6)$? That is, what's y when x is 6?

x	y
6	?


Regression example 1/2

- We have the data from 10 thousand houses in the U.S.

$m=10000$ {


City	Rooms	SQM	Price
Los Angeles	3	130	420000
Los Angeles	2	60	380000
...
Albuquerque	2	140	220000
Albuquerque	3	150	250000

- **Goal:** learn a function (model) that can infer the *Price* given all the other variables, for houses not in the 10 thousand dataset:

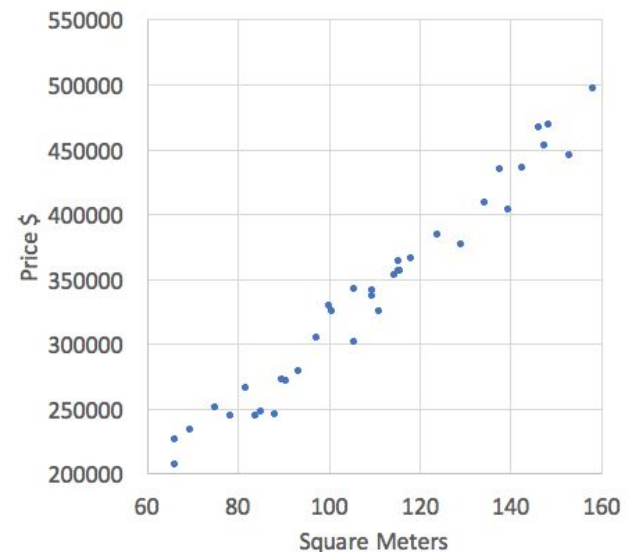


Regression example 2/2

- The houses' prices example is an example of **regression** task:
 - Among the variables there is the output variable (Price)
 - The output variable is the *ground truth*, because it's the actual price of the house
 - The output variable we want to predict is a numerical value



- To visualize things better we will have examples with only two variables (one used as input one as output)
- All the process can be generalized for n variables



Linear regression

- We defined regression as the general task of predicting the real value of a variable using the other variables as input
- The model we want to learn is a function f from n real values x_i to one real value y :

$$y = f(x_1, \dots, x_n)$$

- The learned model is called *hypothesis* function h :

$$y = h(x_1, \dots, x_n)$$


- When **h is a linear function** on the input variables x_i , the regression task is called **linear regression**

Linear function



- A linear function is a function of this form:

$$y = f(x_1, \dots, x_n) = w_1 x_1 + \dots + w_n x_n + b$$


- When $n = 1$ (only one input variable), the plotted function has the **shape of a straight line**. When $n = 2$ it takes the shape of a straight plane etc.
- Straight lines and straight planes are examples of linear functions



**Linear
functions**



**Non-linear
functions**



Linear regression: parameters


- Recall, a linear function is a function of this form:

$$y = f(x_1, \dots, x_n) = w_1 x_1 + \dots + w_n x_n + b$$

- This means that the function that models our data is already defined!
- What's missing? **What are the parameters we need to “learn”?**
 - The x_i variables are the ones **given in input**, so we already have them (e.g., the square meters of the house)
 - The w_i **coefficients are not known!**
 - In a straight line, the coefficient is the **slope** of the linear function!
 - In ML they are also called **weights**: assuming that you have scaled your variables (e.g. between 0 and 1) they tell you **how much a variable contributes** to the final result (the y)
 - The b parameter is **also not known!**
 - In a straight line, this is the **y-intercept** of the line: the point where the line intercept the y axis
 - In ML this is also called **bias term** or **bias weight**

Example: data points

- Let's clarify these new notions with an example
- We have plotted some data points, each point is an observation with two variables: x and y



Example: slope 1/3


- Let's simulate a learning process: we need to learn the slope on the line
- To keep things simple, we focus on the angle only, setting the y-intercept to 0

1. **Let's try with a slope $w = 1$:** this means the line will have a 45° slope
2. Having a y-intercept (b parameters) set to 0, the hypothesis function is:

$$y = h(x) = wx = x$$

3. By plotting the line on the figure we can visually evaluate how much this hypothesis match the actual data

- Is this hypothesis good enough?
- We will see later how to **measure precisely the distance between the actual data and the hypothesis**
- But for sure can do better!




Example: slope 2/3

- Let's simulate a learning process: we need to learn the slope on the line
- To keep things simple, we focus on the angle only, setting the y-intercept to 0

1. **Now let's try with a slope $w = 2$!**
2. Having a y-intercept (b parameters) set to 0, the hypothesis function is:

$$y = h(x) = 2x$$

3. By plotting the line on the figure we can visually evaluate how much this hypothesis match the actual data
 - Is this hypothesis better?
 - The hypothesis **seems worse than before**




Example: slope 3/3

- Let's simulate a learning process: we need to learn the slope on the line
- To keep things simple, we focus on the angle only, setting the y-intercept to 0

1. **Now let's try with a slope $w = 0.5$!**
2. Having a y-intercept (b parameters) set to 0, the hypothesis function is:

$$y = h(x) = \frac{1}{2} x$$

3. By plotting the line on the figure we can visually evaluate how much this hypothesis match the actual data
 - Is this hypothesis better?
 - The hypothesis seems **quite accurate, at least for the slope**
 - We should now learn **what's the best value for the bias term b**



Example: bias term 1/2

- We have now a **good slope but a clearly wrong bias term** (the y intercept parameter)
- Let's try again some reasonable value


1. **Now let's try with a bias $b = 3$!**

2. Having already set the weight to 0.5, the hypothesis function is:

$$y = h(x) = \frac{1}{2}x + 3$$

3. By plotting the line on the figure we can visually evaluate how much this hypothesis match the actual data

- Are the points approximately around the line?
- **It seems slightly above. We should try a lower value for b**



Example: bias term 2/2


- We have now a **good slope but a clearly wrong bias term** (the y intercept parameter)
- Let's try again some reasonable value

1. **Now let's try with a bias $b = 2$!**
2. Having already set the weight to 0.5, the hypothesis function is:

$$y = h(x) = \frac{1}{2}x + 3$$


3. By plotting the line on the figure we can visually evaluate how much this hypothesis match the actual data

- **This hypothesis looks quite right!**



Example: observations

- The machine learning algorithm we have simulated is rather naive:
 - It has no precise way to measure how close we are to the optimal hypothesis
 - It makes almost random guesses to generate new hypothesis
- This means that a good machine learning model should have:
 - **A way to measure how good (or how bad) an hypothesis is**
 - **A smart algorithm that tunes the weights** until the measure of badness is very low (or conversely the measure of goodness is very high)





Loss function

- In ML, the measure used to evaluate the hypothesis is called a **loss function**
 - Intuitively, a **loss** function measures how **bad** the model is with respect to the actual data
- The most common it's the mean squared error, that is the **average squared distance** between the training data and the hypothesis
 - Mean because we want to consider all the m data points in the training set, so we average all the distances
 - Squared to enforce the distance to be non-negative

The diagram illustrates the Mean Squared Error (MSE) formula:
$$\frac{1}{m} \sum_{i=1}^m (h(x_i) - y_i)^2$$
 Annotations include:

- A dashed oval around the $\frac{1}{m} \sum_{i=1}^m$ term, with a dashed line pointing to the text "Arithmetic mean".
- A bracket under the $(h(x_i) - y_i)$ term, with a dashed line pointing to the text "Difference between the y predicted using the hypothesis function and the actual y".
- A dashed arrow pointing from the squared term 2 to the text "We square the difference to enforce non-negativity".

Optimization

- Now that we have a tool to measure the error of an hypothesis function, we want the error to be the smallest possible
- This means the our **objective function** is to find the variables values (weights and bias) that **minimize** the error

$$\text{minimize} : \frac{1}{m} \sum_{i=1}^m (h(x_i) - y_i)^2$$

- Using the *argmin* notation and replacing $h(x)$ with the linear function:

$$\operatorname{argmin}_{w,b} \frac{1}{m} \sum_{i=1}^m ((wx_i + b) - y_i)^2$$




Optimization algorithms


- Most popular and effective optimization algorithm in ML is **gradient descent**
- It **iterates** through 4 steps until the **loss is smaller than a tolerance value**:
 1. Compute the gradient (that is the derivative or slope) of the loss function
 2. The **sign of the gradient** tells us if the loss increase (positive) or decrease (negative) moving to the right. Recall that we want to reach the minimum:
 - a. If it increase, we should move the weight to the left (subtract a small value)
 - b. If it decrease, we should move to the weight to the right (add a small value)
 3. Update the hypothesis with the new weights
 4. Compute again the loss

Gradient descent example 1/6

- For simplicity, we consider only one parameter w to be learned
- The loss function with respect to w it's a convex function, thus it has a global minimum



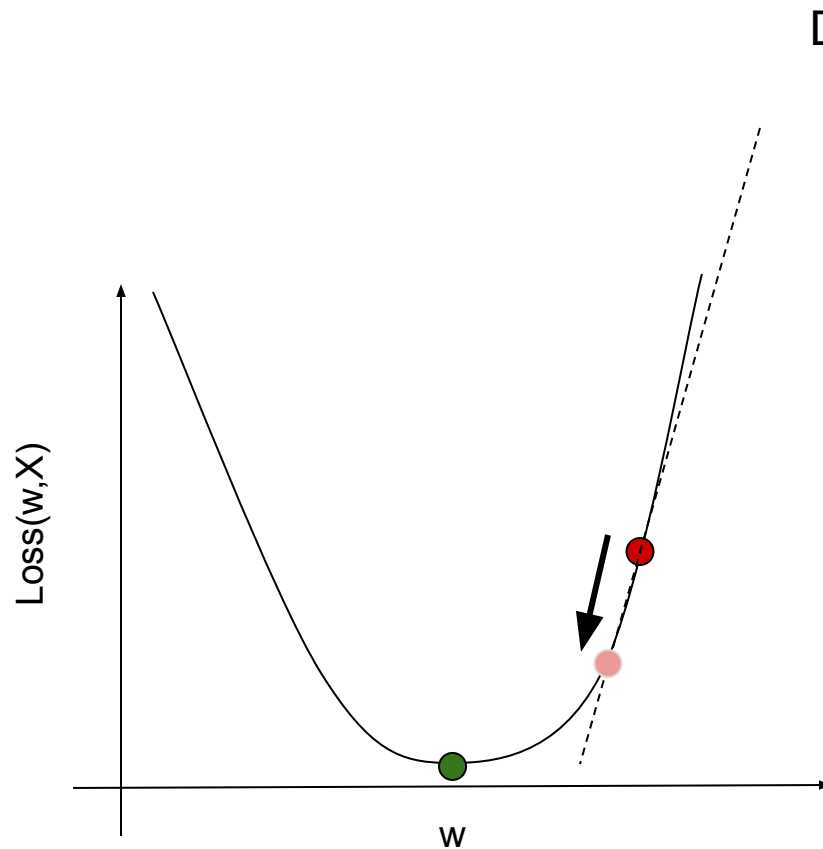
Gradient descent example 2/6



Derivate (gradient) of
the loss function for
the current w value

- The gradient is positive (increasing)
- In order to move toward the minimum, we should subtract a small value from w


Gradient descent example 3/6




Derivate (gradient) of
the loss function for
the current w value

- The gradient is still positive (increasing)
- In order to move toward the minimum, we should subtract a small value from w

Gradient descent example 4/6




Gradient descent example 5/6



- The gradient is still positive (increasing)
- In order to move toward the minimum, we should subtract a small value from w

Gradient descent example 6/6



- The gradient **now is negative** (decreasing)
- In order to move toward the minimum, we should **add a small value to w**
- However, we could decide that the loss is **small enough** and **stop our descent iterations**

Testing set


- **Important:** the loss function is made to evaluate the hypothesis **on the training set!** It measures the distance between the training data and the hypothesis
- If the loss is 0, it means that the hypothesis is **perfectly fitting** the training data
- Yet, this hypothesis could be inaccurate **with unseen data**: data that is not in the training set
 - This phenomenon is called **overfitting**: the model is extremely good (overfitted) on the training data but unable to generalize on new data
- In order to actually test our model, we must use a set of data that the learning algorithm has never seen, that is the **testing set**

Train/test split

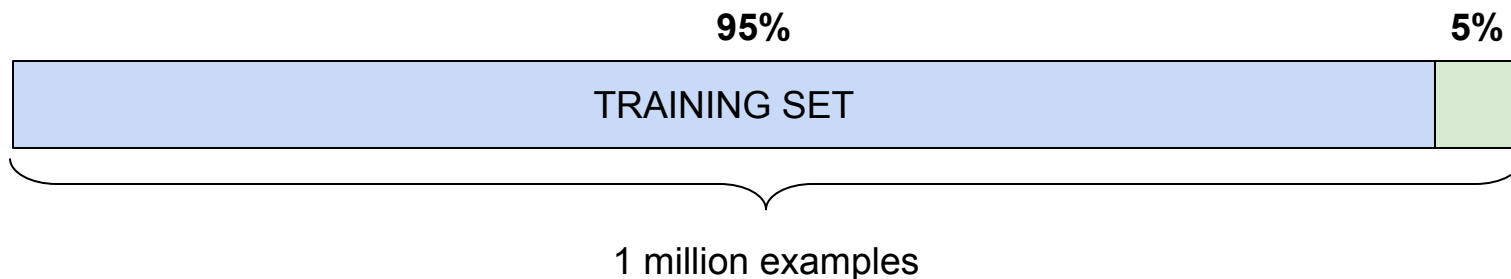
- Usually in a data mining process the data **is not already** splitted into training and testing set
- In splitting between training and testing set you should consider the following:
 - The more data you have in the training the better will perform the learned model
 - The more data you have in the testing the better will be the estimation of accuracy (e.g., if you guess just one example you will have 100% accuracy, but it's not a significant estimation)
- **IMPORTANT:** both training and testing set **must come from the same distribution**. For example:
 - You can't train on the houses' prices in L.A. and test on the houses in N.Y.
 - You can't train on pictures taken with the smartphone and test it on pictures taken from Google Images

Train/test split ratio

- The rule-of-thumb is to split **80/20** following the Pareto principle: 80 for training and 20 for testing



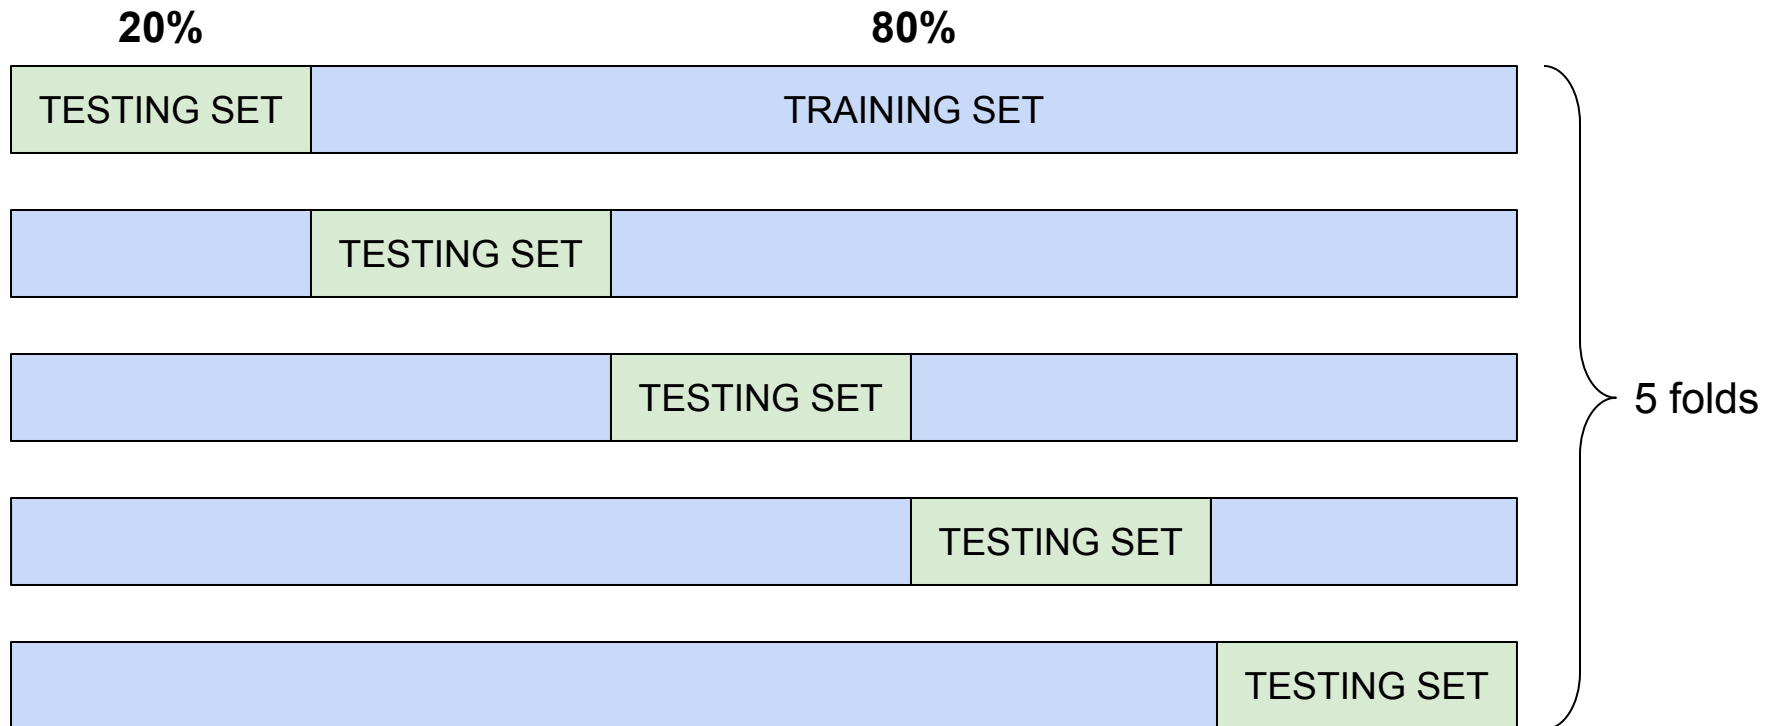
- However, **if you have many examples**, you may decrease the size of testing set to have a more accurate model with a still significant testing set



- What if we want to use **all the data** both for training **and** testing?

Cross-Validation (CV)

- The idea of cross-validation is to rotate over testing/training set split, so that each example will be used for training and for testing
- The fraction of the testing set will determine the number of rotations, also called “folds”. For example, using a testing set of 20%, that is $\frac{1}{5}$, we will have a 5-folds cross validation:





Cross-Validation: limit scenarios

- A **2-fold** cross-validation splits the data 50/50:
 - 1st fold: the first half is used for testing, the other for training
 - 2nd fold: the first half is used for training, the other for testing
- A **m-fold** cross-validation (m is the number of examples in the whole dataset) splits the data $1/m$:
 - 1st fold: first example is used for testing, all the rest ($m-1$ examples) is used for training
 - 2nd fold: second example is used for testing, the first example and all the other examples from the third are used for training
 - ...
 - m -th fold: first $m-1$ examples are used for training, last example is used for testing

m -fold a.k.a. **Leave-One-Out Cross-Validation**

- Note on CV: **the error of the model will be the average of all the errors**: in the m -fold CV it will be the average of m error values

What about non-linear data?


- So far we have seen a regression method that only works if the function from the input variables to the output variable is linear. E.g., with only one input variable x :

$$y = f(x) = wx + b$$


- What if the function is non-linear? The associated expression would be like:

$$y = f(x) = w_1x + w_2x^2 + w_3x^3 \dots$$

with a different weight w for each exponential



**Non-linear
functions**



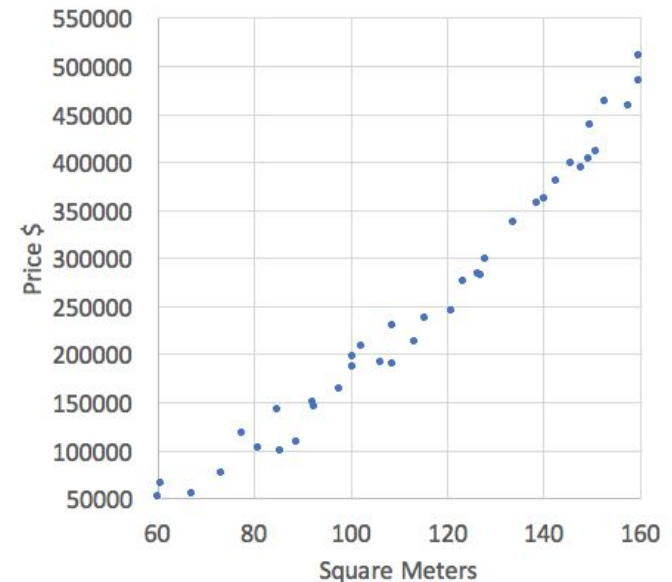
- A simple trick to obtain non-linearity is to artificially generate new features:
 - We start from a feature x_1
 - We can compute the feature x_1^2 and add this new feature to our data
 - We then compute the feature x_1^3 and add it to the data
 - And so on...

Example: house prices


- Let's suppose that the price is a non-linear function of the size in square meters
- In this example, the function that relates the price to the size is:

$$\text{Price} = 0 \cdot \text{Size} + 20 \cdot \text{Size}^2 + 0$$

- We kept the zeros to show all the coefficients and variables
- We can model the above function by augmenting our dataset as follows:



Size (sqm)	Price
60	63193.69799
65.04312529	70911.59045
68.8814367	80082.32417
69.35554098	80602.04965
69.73350476	114729.8332
74.53091304	131513.3546



Size (sqm)	Size ²	Price
60	3600	63193.69799
65.04312529	3887.220733	70911.59045
68.8814367	4436.82399	80082.32417
69.35554098	4764.013352	80602.04965
69.73350476	5225.380597	114729.8332
74.53091304	5321.441446	131513.3546



Limits and advanced methods

- There are limits to the previous “trick”:
 - We have to **manually** add features before training
 - We don’t know where to **stop**
 - to which degree of the polynomial? x^3 ? x^4 ? ... x^{20} ?
 - How about products of features like x_1x_2 or $x_1^2x_2^3x_2^2$?
 - There can be so many features that this would make dimensionality (number of features) to explode
- There are **advanced methods** for regression that do not need to manually add features but have different ways of coping with non-linearity

Classification

- The classification problem ask to predict a categorical variable's value from the values of other variables
- Easiest example is with two variables x and y:
 - x is the variable in input
 - y is the variable we want to predict

x	y
1	A
2	A
3	B
4	B
5	A

Training data: x is the input data, y is the label data.

Learning

Learning task: **what's the mapping from x to y?**

Testing

Try to “learn” from the training data an hypothesis function h

What is $h(6)$? That is, what's y when x is 6?


x	y
6	?

Classification example 1/2

- We have again the data from 10 thousand houses in the U.S.


City	Rooms	SQM	Price
Los Angeles	3	130	420000
Los Angeles	2	60	380000
...
Albuquerque	2	140	220000
Albuquerque	3	150	250000

- Goal: learn a function (model) that can infer the City given all the other variables, for houses not in the 10 thousand dataset:




Classification example 2/2

- The houses' city can be an example of **classification** task:
 - Among the variables we choose the output variable (City)
 - The output variable is the *ground truth*, because it's the actual city of the house
 - The output variable we want to predict is a categorical value

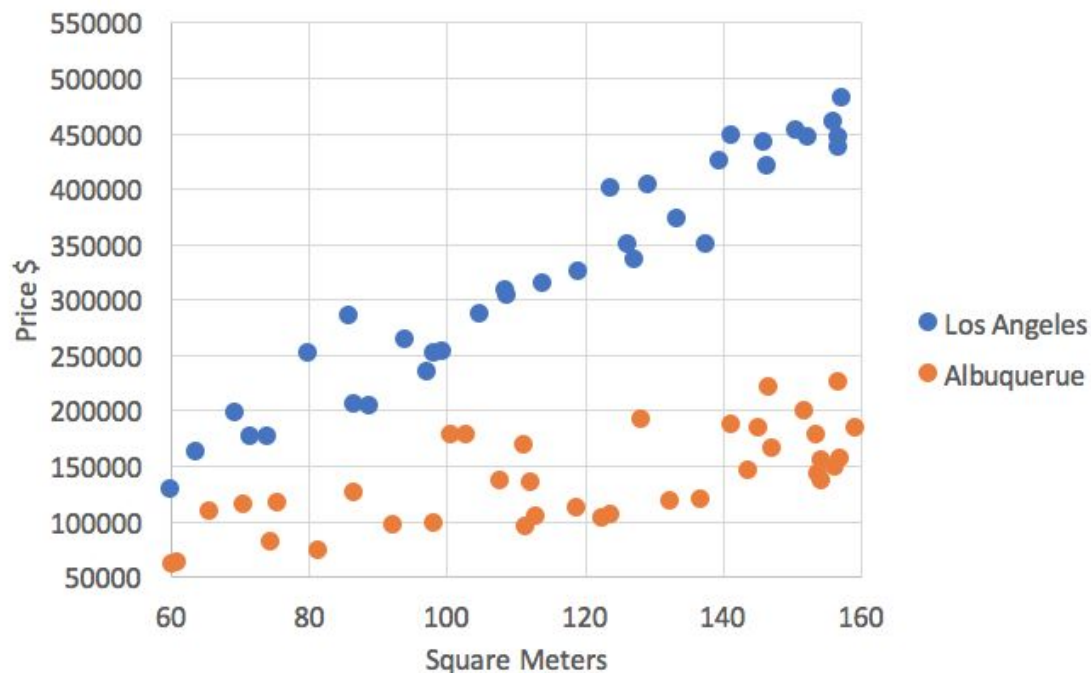


- To visualize things better we will have examples with two variables (no rooms) as input and one as output
- The categorical output will be visualized using a **different color for each category**



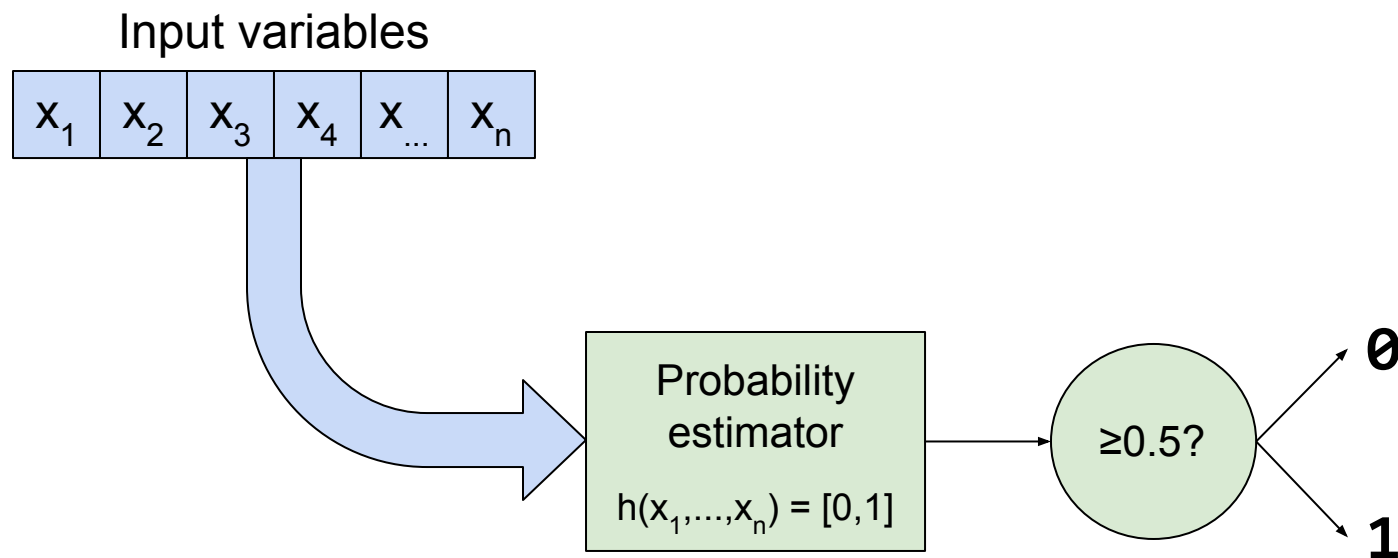
Binary classification

- The binary classification is the task of classifying between two classes.
For example:
 - Detecting if an email is spam or not
 - Deciding if an image has a cat in it
 - Predicting if a currency will go up or down
 - Classifying if a house is in L.A. or Albuquerque




Binary classifier prediction

- A binary classifier prediction works as following:
 - Estimate a score of probability for the first class
 - Output a binary value:
 - 1 if the probability score is greater than or equal to 0.5
 - 0 if the probability score is less than 0.5



Separating the space

- Visually speaking, it is very easy to linearly separate (i.e. with a straight line) the space between Los Angeles houses and Albuquerque houses:



- But how could we **learn** a function that separate the space automatically?




Linear regression for classification?

- In the case of classification, our hypothesis is a function that serves as a probability estimator
 - It takes the variables in input (e.g., house size and price)
 - It must output **a value between 0 and 1**
- Can we use **linear regression**? After all ...
 - ... the output value is actually numerical: 0 or 1
 - ... the line that separates our examples is a straight line
- **Problem:** we don't have control over unseen data, so we could have values bigger than 1, or lower than 0
- **Idea:**
 - We use a linear function to combine all the input variables into a value
 - We apply a function to constrain every possible value into a range of $[0,1]$

Logistic (sigmoid) function

- The ***sigmoid*** function is a special case of the **logistic function**
- Sigmoid functions have domain of **all real numbers**, with return value monotonically increasing **from 0 to 1**



Logistic Regression

- Despite the name, the logistic regression is actually a **classification** method
- Recall the linear regression hypothesis was a linear combination of the input variables, with one weight for each variable:

$$h_{\text{linear}}(x_1, \dots, x_n) = w_1 x_1 + \dots + w_n x_n + b$$

- Its hypothesis function is simply the sigmoid function applied to the linear combination:

$$h_{\text{logistic}}(x_1, \dots, x_n) = \text{sig}(w_1 x_1 + \dots + w_n x_n + b)$$



Logistic Regression: loss function

- Recall that we need a loss function to **tune the weights towards a minimum loss**
- We have a training input x . We consider two cases:
 - **When the true y is 0:** the hypothesis should output 0, anything more than 0 is a loss! So we can use simply the hypothesized score


$$\text{loss}(x) = h_{\text{logistic}}(x)$$

- **When the true y is 1:** the hypothesis should output 1, anything less than 1 is a loss! So we can use the difference between 1 and the hypothesized score

$$\text{loss}(x) = 1 - h_{\text{logistic}}(x)$$

Problem: Overfitting 1/2


- Using the **same trick** of introducing artificial features using polynomials of different degrees (e.g. $x_1^2 x_2^3 x_2^2 \dots$), we could train a very complex non-linear function:



- Is this a good thing?** For sure it's perfect for the training set ... what about **unseen data**?



Problem: Overfitting 2/2

- It turns out it is actually not good: the model is **overly fitted** on the training set the unseen data could fall randomly on one side or the other
- The model should be more robust and **generalized** to work also on unseen data (such as the testing set)



Solution: Regularization

- Regularization is a very effective technique to prevent overfitting
 - It's a modification of the loss function
 - We add an additional value to the loss function that increases as the value of features weights (w) increase
- What is the effect of regularization?
 - It prevents the features weights (w) to be very large
 - Visually speaking, it will “smooth” the decision boundary line:



Multi-class classification

- Binary classification **can be generalized** to multi-class classification
- There are **two ways** to adapt binary classifiers to a **K-classes** classification problem:
 - **OVO - One vs. One**: Each possible pair of classes is taken in consideration to train a different classifier. This approach will **need to train $K*(K-1)/2$ classifiers!**
 - **OVA - One vs. All** (more common): for each class, a different classifiers is trained to distinguish between that class and all the other classes merged together. This approach will **need to train only K classifiers**



References

[Stanford Lectures on Machine Learning](#)

Andrew Ng